

Identifying Drivers of Outcomes: Linear Models

Luis Francisco Gómez López

FAEDIS

2025-05-05



UNIVERSIDAD MILITAR
NUEVA GRANADA

Table of contents I

- 1 Please Read Me
- 2 Purpose
- 3 Amusement park survey
- 4 Acknowledgments

- This presentation is based on (Chapman and Feit 2019, chap. 7)

- Apply linear modeling to understand a response variable and make predictions of forecasts

- **weekend**: whether the visit was on a weekend
- **num.child**: number of children in the visit
- **distance**: how far the customer traveled to the park in miles
- **rides**: satisfaction with rides using a scale $[0, 100]$
- **games**: satisfaction with games using a scale $[0, 100]$
- **wait**: satisfaction with waiting times using a scale $[0, 100]$
- **clean**: satisfaction with cleanliness using a scale $[0, 100]$
- **overall**: overall satisfaction rating using a scale $[0, 100]$

• Import data

```
amusement_park <- read_csv("http://goo.gl/HKn174")
amusement_park |> head(n = 5)
```

```
# A tibble: 5 x 8
  weekend num.child distance rides games wait clean overall
  <chr>      <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>
1 yes          0    115.    87    73    60    89     47
2 yes          2    27.0    87    78    76    87     65
3 no           1    63.3    85    80    70    88     61
4 yes          0    25.9    88    72    66    89     37
5 no           4    54.7    84    87    74    87     68
```

• Transform data

```
amusement_park <- amusement_park |>
  mutate(weekend = factor(x = weekend,
                          labels = c('no', 'yes'),
                          ordered = FALSE),
         num.child = as.integer(num.child),
         # logarithmic transform
         logdist = log(distance, base = exp(x = 1)))
amusement_park |> head(n = 5)
```

A tibble: 5 x 9

	weekend <fct>	num.child <int>	distance <dbl>	rides <dbl>	games <dbl>	wait <dbl>	clean <dbl>	overall <dbl>	logdist <dbl>
1	yes	0	115.	87	73	60	89	47	4.74
2	yes	2	27.0	87	78	76	87	65	3.30
3	no	1	63.3	85	80	70	88	61	4.15
4	yes	0	25.9	88	72	66	89	37	3.25
5	no	4	54.7	84	87	74	87	68	4.00

- Summarize data
 - Ups the table is really big!!! Try it in your console to see the complete table

```
amusement_park |> skim()
```


- Correlation matrices

- Pearson correlation coefficients for samples in a tibble

```
correlation_matrix <- amusement_park |>
  select(num.child, rides:logdist) |>
  corrr::correlate()
correlation_matrix
```


A tibble: 7 x 8

term <chr>	num.child <dbl>	rides <dbl>	games <dbl>	wait <dbl>	clean <dbl>	overall <dbl>	logdist <dbl>
1 num.child	NA	-0.0403	0.00466	-0.0210	-0.0135	0.319	-0.00459
2 rides	-0.0403	NA	0.455	0.314	0.790	0.586	-0.0110
3 games	0.00466	0.455	NA	0.299	0.517	0.437	0.00187
4 wait	-0.0210	0.314	0.299	NA	0.368	0.573	0.0175
5 clean	-0.0135	0.790	0.517	0.368	NA	0.639	0.0221
6 overall	0.319	0.586	0.437	0.573	0.639	NA	0.0763
7 logdist	-0.00459	-0.0110	0.00187	0.0175	0.0221	0.0763	NA

- Correlation matrices


- Pearson correlation coefficients for samples in a tibble

```
correlation_matrix |> autoplot(triangular = "lower")
```




• Bivariate Association: the base R way

```
plot(overall-rides, data=amusement_park,
     xlab="Satisfaction with Rides", ylab="Overall Satisfaction")
abline(reg = lm(formula = overall-rides, data = amusement_park),
       col = 'blue')
```




• Bivariate Association: the tidyverse way

```
amusement_park |> ggplot(aes(x = rides, y = overall)) +
  geom_point() +
  geom_smooth(method = 'lm',
              color = 'blue',
              se = FALSE) +
  labs(x = "Satisfaction with Rides",
       y = "Overall Satisfaction")
```



Linear Model with a Single Predictor



- Linear Model with a Single Predictor

$overall_i = \beta_0 + \beta_1 rides_i + \epsilon_i$ where $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$ and $i = 1, \dots, 500$

$\widehat{overall}_i = \hat{\beta}_0 + \hat{\beta}_1 rides_i$ and $\hat{\sigma}^2$ where $i = 1, \dots, 500$

$overall_i - \widehat{overall}_i = \hat{\epsilon}_i$ where $i = 1, \dots, 500$

```
model1 <- lm(formula = overall ~ rides, data = amusement_park)
model1
```

Call:

```
lm(formula = overall ~ rides, data = amusement_park)
```

Coefficients:

(Intercept)	rides
-94.962	1.703

• Linear Model with a Single Predictor

```
ls.str(model1)
```

```
assign : int [1:2] 0 1
call : language lm(formula = overall ~ rides, data = amusement_park)
coefficients : Named num [1:2] -95 1.7
df.residual : int 498
effects : Named num [1:500] -1146.2 -207.9 11.5 -17.9 20.3 ...
fitted.values : Named num [1:500] 53.2 53.2 49.8 54.9 48.1 ...
model : 'data.frame': 500 obs. of 2 variables:
 $ overall: num 47 65 61 37 68 27 40 30 58 36 ...
 $ rides : num 87 87 85 88 84 81 77 82 90 88 ...
qr : List of 5
 $ qr : num [1:500, 1:2] -22.3607 0.0447 0.0447 0.0447 0.0447 ...
 $ qraux: num [1:2] 1.04 1.01
 $ pivot: int [1:2] 1 2
 $ tol : num 1e-07
 $ rank : int 2
rank : int 2
residuals : Named num [1:500] -6.22 11.78 11.18 -17.93 19.89 ...
terms : Classes 'terms', 'formula' language overall ~ rides
xlevels : Named list()
```

• Linear Model with a Single Predictor

```
summary(model1)
```

Call:

```
lm(formula = overall ~ rides, data = amusement_park)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-33.597	-10.048	0.425	8.694	34.699

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-94.9622	9.0790	-10.46	<2e-16 ***
rides	1.7033	0.1055	16.14	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 12.88 on 498 degrees of freedom

Multiple R-squared: 0.3434, Adjusted R-squared: 0.3421

F-statistic: 260.4 on 1 and 498 DF, p-value: < 2.2e-16

- Linear Model with a Single Predictor

```
model1$coefficients
```

```
(Intercept)      rides
-94.962246      1.703285
```

```
# Make some predictions
# We want to forecast the overall satisfaction rating
# if the satisfaction with rides is 95
-94.962246 + 1.703285*95
```

```
[1] 66.84983
```

- Linear Model with a Single Predictor
 - Std. Error column
 - Indicates uncertainty in the coefficient estimate
 - We can build a confidence interval

```
summary(model1)$coefficients[, 2]
```

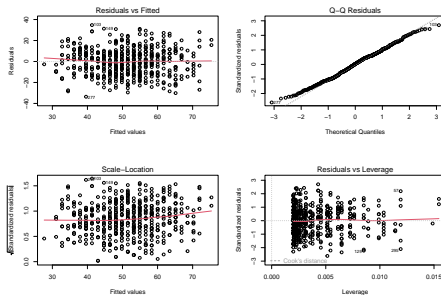
```
(Intercept)      rides
  9.0790049    0.1055462
```

```
confint(model1, level = 0.95)
```

```
                2.5 %      97.5 %
(Intercept) -112.800120 -77.124371
rides        1.495915   1.910656
```

- Linear Model with a Single Predictor

```
par(mfrow=c(2,2))
plot(model1)
```



```
par(mfrow=c(1,1))
```

- Linear Model with a Single Predictor
 - **Linearity:** plot (1, 1)
 - Reference line should be flat and horizontal
 - **Normality of residuals:** plot (1, 2)
 - Dots should fall along the line
 - **Homogeneity of variance:** plot (2, 1)
 - Reference line should be flat and horizontal
 - **Influential observations:** plot (2, 2)
 - Points should be inside the contour lines

- Linear Model with Multiple Predictors

$$\begin{aligned} overall_i &= \beta_0 + \beta_1 rides_i + \beta_2 games_i \\ &\quad + \beta_3 wait_i + \beta_4 clean_i + \epsilon_i \\ \text{where } \epsilon_i &\sim \mathcal{N}(0, \sigma^2) \text{ and } i = 1, \dots, 500 \end{aligned}$$

```
model2 <- lm(formula = overall ~ rides + games + wait + clean,
             data = amusement_park)
model2
```

Call:

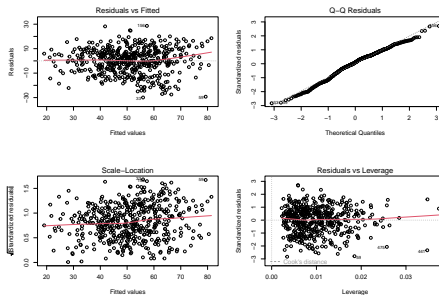
```
lm(formula = overall ~ rides + games + wait + clean, data = amusement_park)
```

Coefficients:

(Intercept)	rides	games	wait	clean
-131.4092	0.5291	0.1533	0.5533	0.9842

• Linear Model with Multiple Predictors

```
par(mfrow=c(2,2))
plot(model2)
```



```
par(mfrow=c(1,1))
```

• Linear Model with Multiple Predictors

```
summary(model2)
```

Call:

```
lm(formula = overall ~ rides + games + wait + clean, data = amusement_park)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-29.944	-6.841	1.072	7.167	28.618

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-131.40919	8.33377	-15.768	< 2e-16 ***
rides	0.52908	0.14207	3.724	0.000219 ***
games	0.15334	0.06908	2.220	0.026903 *
wait	0.55333	0.04781	11.573	< 2e-16 ***
clean	0.98421	0.15987	6.156	1.54e-09 ***

 Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 10.59 on 495 degrees of freedom

Multiple R-squared: 0.5586, Adjusted R-squared: 0.5551

F-statistic: 156.6 on 4 and 495 DF, p-value: < 2.2e-16

- Linear Model with Multiple Predictors

$$H_0 : \beta_1 = 0$$

$$H_1 : \beta_1 \neq 0$$

$$t_{rides} = \frac{\hat{\beta}_1 - \beta_1}{\sqrt{Var(\hat{\beta}_1)}} = \frac{0.529078 - 0}{0.14207176} = 3.724019$$


```
model2$coefficients
```

(Intercept)	rides	games	wait	clean
-131.4091939	0.5290780	0.1533361	0.5533264	0.9842126

```
# Calculate the variance-covariance matrix, extract
# the diagonal and calculate the standard deviation of
# the parameters
model2 |> vcov() |> diag() |> sqrt()
```

(Intercept)	rides	games	wait	clean
8.33376643	0.14207176	0.06908486	0.04781282	0.15986712

Linear Model with Multiple Predictors




• Linear Model with Multiple Predictors

```
confint(model2, level = 0.95)
```

	2.5 %	97.5 %
(Intercept)	-147.78311147	-115.0352764
rides	0.24993998	0.8082161
games	0.01760038	0.2890718
wait	0.45938535	0.6472675
clean	0.67011082	1.2983144

• Linear Model with Multiple Predictors

```
library(coefplot) # Remember to install the package if it is not installed
coefplot(model = model2,
  # The intercept is relatively large: -131.4092
  intercept = FALSE,
  ylab="Rating of Feature",
  xlab="Association with Overall Satisfaction",
  lwdOuter = 1.5)
```



- Comparing models

```
summary(model1)$r.squared
```

```
[1] 0.3433799
```

```
summary(model2)$r.squared
```

```
[1] 0.558621
```

```
summary(model1)$adj.r.squared
```


```
[1] 0.3420614
```

```
summary(model2)$adj.r.squared
```

```
[1] 0.5550543
```

- Comparing models
 - Base R way

```
plot(x = amusement_park$overall, y = fitted(model1),
     col = "red", xlim = c(0,100), ylim = c(0,100),
     xlab = "Actual Overall Satisfaction",
     ylab = "Fitted Overall Satisfaction")
points(x = amusement_park$overall, y = fitted(model2),
       col = "blue")
legend(x = "topleft", legend = c("Model 1", "Model 2"), col = c("red", "blue"), pch = 1)
```



- Comparing models

- Tidymodels and tidyverse way: Prepare data

```
model1_augment <- augment(x = model1) |> mutate(model = "Model 1")
model2_augment <- augment(x = model2) |> mutate(model = "Model 2")
models_performance <- model1_augment |> bind_rows(model2_augment)

models_performance |> glimpse()
```

```
Rows: 1,000
Columns: 12
$ overall    <dbl> 47, 65, 61, 37, 68, 27, 40, 30, 58, 36, 71, 48, 75, 46, 59,~
$ rides      <dbl> 87, 87, 85, 88, 84, 81, 77, 82, 90, 88, 93, 79, 94, 81, 86,~
$ .fitted    <dbl> 53.22359, 53.22359, 49.81702, 54.92688, 48.11373, 43.00388,~
$ .resid     <dbl> -6.2235914, 11.7764086, 11.1829795, -17.9268769, 19.8862650~
$ .hat       <dbl> 0.002089430, 0.002089430, 0.002048063, 0.002311576, 0.00222~
$ .sigma     <dbl> 12.88964, 12.88182, 12.88289, 12.86751, 12.86171, 12.87260,~
$ .cooks     <dbl> 2.449537e-04, 8.770564e-04, 7.751689e-04, 2.249493e-03, 2.6~
$ .std.resid <dbl> -0.48371422, 0.91529407, 0.86915315, -1.39348008, 1.5457218~
$ model      <chr> "Model 1", "Model 1", "Model 1", "Model 1", "Model 1", "Mod~
$ games      <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
$ wait       <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
$ clean      <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
```

- Comparing models
 - Tidymodels and tidyverse way: Visualize

```
models_performance |>
  ggplot() +
  geom_point(aes(x = overall, y = .fitted,
                 color = model)) +
  labs(x = "Actual Overall Satisfaction",
       y = "Fitted Overall Satisfaction")
```



- Comparing models
 - Analysis of variance (anova) for nested models¹

```
anova_lm <- anova(model1, model2, test = "F")
anova_lm
```

Analysis of Variance Table

Model 1: overall ~ rides

Model 2: overall ~ rides + games + wait + clean

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	498	82612				
2	495	55532	3	27080	80.463	< 2.2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


¹This statistical analysis only make sense for nested models that are fitted with the same data where the convention is to include the models from smallest to largest. See `?anova_lm`

• Comparing models

$$H_0 : \beta_0 = \beta_1 = \beta_2 = \beta_3 = \beta_4 = 0$$

$$H_1 : \text{At least one } \beta_j \neq 0 \text{ for } j = 0, 1, 2, 3, 4$$

$$F = \frac{\frac{RSS_1 - RSS_2}{p_2 - p_1}}{\frac{RSS_2}{n - p_2}} = \frac{\frac{82611.81 - 55531.53}{5 - 2}}{\frac{55531.53}{500 - 5}} = 80.46323$$



- Predictions

$$\widehat{overall}_j = \hat{\beta}_0 + \hat{\beta}_1 rides_j + \hat{\beta}_2 games_j + \hat{\beta}_3 wait_j + \hat{\beta}_4 clean_j$$

```
coef(model2) |> enframe(name = "coef")
```

```
# A tibble: 5 x 2
  coef      value
  <chr>    <dbl>
1 (Intercept) -131.
2 rides        0.529
3 games        0.153
4 wait         0.553
5 clean        0.984
```

- Predictions

- Manual

```
(coef(model2)["(Intercept)"]*1 + coef(model2)["rides"]*30 + coef(model2)["games"]*10 +  
  coef(model2)["wait"]*57 + coef(model2)["clean"]*90) |>  
  unname()
```

```
[1] 6.11525
```

- Predictions

- Matrix multiplication

```
coef(model2) %*% c(1, 30, 10, 57, 90)
```

```
      [,1]  
[1,] 6.11525
```

- Predictions
 - predict

```
# New data
new_data <- tibble(rides = c(30, 70),
                    games = c(10, 80),
                    wait = c(57, 60),
                    clean = c(90, 93))

# Result
predict(object = model2, newdata = new_data) |>
  enframe(name = "observation", value = "overall_pred") |>
  bind_cols(new_data)
```

```
# A tibble: 2 x 6
  observation overall_pred rides games wait clean
  <chr>          <dbl> <dbl> <dbl> <dbl> <dbl>
1 1            6.12    30    10    57    90
2 2           42.6     70    80    60    93
```

- Standardizing the predictors

- Compare the effect that different predictor variables have on a response variable
- It must be interpreted in terms of standard deviations
 - One standard deviation in x variable is associated with a standard deviation increase or decrease depending on the value of the estimated parameter

```
amusement_park_std <- amusement_park |>
  select(-distance) |>
  mutate(across(rides:logdist,
    .fns = ~ scale(x = .x,
      center = TRUE,
      scale = TRUE)[,1]))
amusement_park_std |> head()
```

```
# A tibble: 6 x 8
  weekend_num.child rides    games    wait    clean overall logdist
  <fct>          <int>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 yes              0  0.211 -0.698 -0.919    0.215  -0.268  1.79
2 yes              2  0.211 -0.0820  0.567   -0.176   0.865  0.323
3 no               1 -0.155  0.164  0.00966  0.0199   0.614  1.19
4 yes              0  0.394 -0.821 -0.362    0.215  -0.898  0.280
5 no               4 -0.338  1.03  0.381   -0.176   1.05  1.04
6 no               5 -0.887  0.0411 -2.03   -1.74   -1.53  0.145
```

• Standardizing the predictors

```
model2_std <- lm(formula = overall ~ rides + games + wait + clean,
                 data = amusement_park_std)
summary(model2_std)
```

Call:
lm(formula = overall ~ rides + games + wait + clean, data = amusement_park_std)

Residuals:

	Min	1Q	Median	3Q	Max
	-1.88578	-0.43082	0.06749	0.45136	1.80231

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.607e-16	2.983e-02	0.000	1.000000
rides	1.820e-01	4.888e-02	3.724	0.000219 ***
games	7.844e-02	3.534e-02	2.220	0.026903 *
wait	3.753e-01	3.243e-02	11.573	< 2e-16 ***
clean	3.170e-01	5.150e-02	6.156	1.54e-09 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.667 on 495 degrees of freedom
Multiple R-squared: 0.5586, Adjusted R-squared: 0.5551
F-statistic: 156.6 on 4 and 495 DF, p-value: < 2.2e-16

• Using factors as predictors

```
model3 <- lm(formula = overall ~ rides + games + wait + clean + weekend + logdist + num.child,
             data = amusement_park_std)
tidy(model3)
```

```
# A tibble: 8 x 5
  term      estimate std.error statistic  p.value
  <chr>      <dbl>    <dbl>    <dbl>    <dbl>
1 (Intercept) -0.373    0.0465   -8.01 8.41e-15
2 rides        0.213    0.0420    5.07 5.57e- 7
3 games        0.0707   0.0303    2.34 1.99e- 2
4 wait         0.381    0.0278   13.7 1.45e-36
5 clean        0.297    0.0441    6.72 4.89e-11
6 weekendyes   -0.0459   0.0514   -0.893 3.73e- 1
7 logdist      0.0647   0.0257    2.52 1.22e- 2
8 num.child    0.227    0.0171   13.3 1.37e-34
```

```
glance(model3)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic  p.value    df logLik   AIC   BIC
  <dbl>      <dbl>    <dbl>    <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
1  0.679      0.674 0.571    148. 5.97e-117     7 -425.  868.  906.
# i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

- Using factors as predictors

- Overall satisfaction is about the same regardless the number of children

```
amusement_park_std <- amusement_park_std |>
  mutate(num.child.factor = factor(num.child))
model4 <- lm(formula = overall ~ rides + games + wait + clean + weekend + logdist + num.child.factor,
             data = amusement_park_std)
tidy(model4) |> slice(1, 2, 8:12)
```

```
# A tibble: 7 x 5
```

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	-0.691	0.0449	-15.4	7.00e-44
2	rides	0.223	0.0354	6.30	6.61e-10
3	num.child.factor1	1.02	0.0713	14.3	8.96e-39
4	num.child.factor2	1.04	0.0564	18.4	8.77e-58
5	num.child.factor3	0.980	0.0702	14.0	1.75e-37
6	num.child.factor4	0.932	0.0803	11.6	1.22e-27
7	num.child.factor5	1.00	0.104	9.66	2.50e-20

```
glance(model4)
```

```
# A tibble: 1 x 12
```

	r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.775	0.770	0.480	153.	2.68e-150	11	-336.	698.	753.

```
# i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```


- Using factors as predictors
 - Preparing data

```
amusement_park_std <- amusement_park_std |>
  mutate(has.child = factor(x = num.child > 0, labels = c("No", "Yes")))
model5 <- lm(formula = overall ~ rides + games + wait + clean + logdist + has.child,
  data = amusement_park_std)
tidy(model5) |> slice(1, 2, 7)
```

A tibble: 3 x 5

	term <chr>	estimate <dbl>	std.error <dbl>	statistic <dbl>	p.value <dbl>
1	(Intercept)	-0.702	0.0391	-18.0	6.68e-56
2	rides	0.223	0.0351	6.34	5.12e-10
3	has.childYes	1.01	0.0468	21.5	1.08e-72

```
glance(model5)
```

A tibble: 1 x 12

	r.squared <dbl>	adj.r.squared <dbl>	sigma <dbl>	statistic <dbl>	p.value <dbl>	df <dbl>	logLik <dbl>	AIC <dbl>	BIC <dbl>
1	0.774	0.771	0.478	282.	1.03e-155	6	-337.	690.	724.

i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>

- Using factors as predictors

- Maybe having children and the visits on weekends are important for the scores so an interaction will be useful

```
model6 <- lm(formula = overall ~ rides + games + wait + clean + weekend + logdist +
             has.child + rides:has.child + games:has.child + wait:has.child +
             clean:has.child + rides:weekend + games:weekend + wait:weekend +
             clean:weekend, data = amusement_park_std)

tidy(model6) |> slice(9:16)
```

```
# A tibble: 8 x 5
  term                estimate std.error statistic  p.value
<chr>                <dbl>    <dbl>    <dbl>    <dbl>
1 rides:has.childYes  0.0578    0.0731    0.792  4.29e- 1
2 games:has.childYes -0.0640    0.0528   -1.21  2.26e- 1
3 wait:has.childYes   0.351     0.0472    7.42  5.21e-13
4 clean:has.childYes -0.00185   0.0797   -0.0233 9.81e- 1
5 rides:weekendyes    0.0618    0.0678    0.912  3.62e- 1
6 games:weekendyes    0.0185    0.0490    0.377  7.06e- 1
7 wait:weekendyes     0.0352    0.0445    0.791  4.29e- 1
8 clean:weekendyes    -0.0273   0.0710   -0.385  7.01e- 1
```

```
glance(model6)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic  p.value    df logLik   AIC   BIC
  <dbl>      <dbl>    <dbl>    <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
1    0.802    0.796  0.452    130.  3.69e-159    15  -304.  643.  714.
# i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```



- Using factors as predictors
 - Only an interaction was significant

```
model7 <- lm(formula = overall ~ rides + games + wait + clean + logdist + has.child +
             wait:has.child, data = amusement_park_std)
tidy(model7)
```

```
# A tibble: 8 x 5
  term                estimate std.error statistic  p.value
  <chr>                <dbl>    <dbl>    <dbl>    <dbl>
1 (Intercept)        -0.693     0.0368    -18.8 6.91e-60
2 rides               0.213     0.0331     6.42 3.24e-10
3 games              0.0487    0.0239     2.03 4.25e- 2
4 wait               0.151     0.0369     4.09 4.98e- 5
5 clean              0.302     0.0349     8.68 5.94e-17
6 logdist            0.0292    0.0203     1.44 1.50e- 1
7 has.childYes       0.998     0.0442    22.6 4.02e-78
8 wait:has.childYes  0.347     0.0438     7.92 1.59e-14
```


```
glance(model7)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic  p.value    df logLik   AIC    BIC
  <dbl>     <dbl> <dbl>    <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
1  0.800     0.797 0.451    280. 2.96e-167  7 -307.  632.  670.
# i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```



- Using factors as predictors
 - Final model

```
library(dotwhisker) # Remember to install the package if it is not installed
tidy(model7) |>
  dwplot(ci = 0.95,
         dot_args = list(size = 2, color = "black"), whisker_args = list(color = "red"),
         vline = geom_vline(xintercept = 0, color = "black", linetype = 2)) +
  labs(x = "Association with Overall Satisfaction", y = "Rating of Feature")
```



● Formula syntax

Formula in R	Statistical Model
$y \sim x$	$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$
$y \sim -1 + x$	$y_i = \beta_1 x_i + \varepsilon_i$
$y \sim x + z$	$y_i = \beta_0 + \beta_1 x_i + \beta_2 z_i + \varepsilon_i$
$y \sim x + z + x:z$	$y_i = \beta_0 + \beta_1 x_i + \beta_2 z_i + \beta_3 x_i z_i + \varepsilon_i$
$y \sim x * z$	$y_i = \beta_0 + \beta_1 x_i + \beta_2 z_i + \beta_3 x_i z_i + \varepsilon_i$
$y \sim (x + z + w)^2$	$y_i = \beta_0 + \beta_1 x_i + \beta_2 z_i + \beta_3 w_i + \beta_4 x_i z_i + \beta_5 x_i w_i + \beta_6 w_i z_i + \varepsilon_i$
$y \sim (x + z + w)^2 - x:z$	$y_i = \beta_0 + \beta_1 x_i + \beta_2 z_i + \beta_3 w_i + \beta_4 x_i w_i + \beta_5 w_i z_i + \varepsilon_i$
$y \sim x + I(x^2)$	$y_i = \beta_0 + \beta_1 x_i + \beta_1 x_i^2 + \varepsilon_i$

● Try the following models using tidy:

```
lm(formula = overall ~ rides, data = amusement_park_std) |> tidy()
lm(formula = overall ~ -1 + rides, data = amusement_park_std) |> tidy()
lm(formula = overall ~ rides + has.child, data = amusement_park_std) |> tidy()
lm(formula = overall ~ rides + has.child + has.child, data = amusement_park_std) |> tidy()
lm(formula = overall ~ (rides + has.child + weekend)^2,
  data = amusement_park_std) |> tidy()
lm(formula = overall ~ (rides + has.child + weekend)^2 - rides:has.child,
  data = amusement_park_std) |> tidy()
lm(formula = overall ~ rides + I(rides^2) - rides:has.child, data = amusement_park_std) |> tidy()
```



- To my family that supports me
- To the taxpayers of Colombia and the **UMNG students** who pay my salary
- To the **Business Science** and **R4DS Online Learning** communities where I learn **R** and **π -thon**
- To the **R Core Team**, the creators of **RStudio IDE**, **Quarto** and the authors and maintainers of the packages **tidyverse**, **skimr**, **tidymodels**, **dotwhisker**, **kableExtra** and **tinytex** for allowing me to access these tools without paying for a license
- To the **Linux kernel community** for allowing me the possibility to use some **Linux distributions** as my main **OS** without paying for a license

References I

Chapman, Chris, and Elea McDonnell Feit. 2019. *R For Marketing Research and Analytics*. 2nd ed. 2019. Use R! Cham: Springer International Publishing : Imprint: Springer.
<https://doi-org.ezproxy.umng.edu.co/10.1007/978-3-030-14316-9>.